

Federated Learning & Data Privacy, 2024-2025

Second Lab - 4 February 2025

Welcome to the second lab session of the Federated Learning & Data Privacy course! In our first lab, we implemented the Federated Averaging (FedAvg) algorithm, writing the client and aggregator classes, and we performed some preliminary experiments.

RECAP OF EXERCISE 3 - The Effect of Local Epochs

Objective: Analyze how the number of local epochs affects the model's performance in a federated learning setting.

Experiment:

- We ran FedAvg for different numbers of local epochs (e.g., 1, 5, 10, 50, 100).
- We recorded the test accuracy for each setting.

Plot:

- We plotted the local epochs on the x-axis and test accuracy on the y-axis.

Analysis:

- Discuss how local epochs influence model accuracy.
- Were you expecting this result?
- How was the data generated and partitioned in TP1? Justify your answer by examining `data/mnist/generate_data.py` and `data/mnist/utls.py`.

NEW EXERCISES FOR TP2

Goal: In this lab, we will analyze the effects of data heterogeneity, implement client sampling, and explore personalization within federated learning frameworks. A bonus exercise offers the opportunity to deploy a federated learning algorithm in a real, distributed network environment.

To get started, clone the TP2 folder from the lab repository.

EXERCISE 4.1 - The Impact of Data Heterogeneity

Objective: Demonstrate that an increase in the number of local epochs can potentially degrade FedAvg's performance under non-IID data distributions.

- **Preliminary question:** What non-IID data distribution means? Provide examples.
- **Pathological Split:** Familiarize yourself with the concept of "pathological split" as explained in [Communication-Efficient Learning of Deep Networks from Decentralized Data \(Section 3\)](https://arxiv.org/abs/1602.05629) (<https://arxiv.org/abs/1602.05629>). The `pathological_non_iid_split` function has been implemented for you in `data/mnist/utls.py`. Review this method and summarize it briefly.
- **Experiments:** Run the `generate_data.py` script with the `--non_iid` flag and set `n_classes_per_client=2` to partition the MNIST dataset in a non-IID fashion.
- **Plot:** Run experiments to observe how varying the number of local epochs (e.g., 1, 5, 10, 50, 100) influences the model's test accuracy under non-IID data distribution. Plot the relationship between

the number of local epochs and test accuracy.

- **Interpretation:** Briefly comment the results. Were these results expected?

EXERCISE 4.2 - Tackling Data Heterogeneity with FedProx

Objective: Understand how the FedProx algorithm addresses the challenges posed by data heterogeneity in federated learning and compare its performance with the FedAvg algorithm.

- **FedProx Overview:** FedProx modifies the local training objective by introducing a proximal term, which aims to reduce local model drift by penalizing significant deviations from the global model. Review the FedProx algorithm [Federated Optimization in Heterogeneous Networks \(Algorithm 2\) \(https://arxiv.org/abs/1812.06127\)](https://arxiv.org/abs/1812.06127) and our implementation of the ProxSGD class in `utils/optim.py`.
 - **Experiments:** To initiate FedProx experiments, run the `train.py` script with `local_optimizer = "prox_sgd"` and set the proximal term coefficient $\mu = 2$.
 - **Plot:** Replicate the plot from Exercise 4.1, this time evaluating FedProx algorithm.
 - **Analysis:** Discuss the observed differences in performance between FedAvg and FedProx. Are there specific configurations (e.g., number of local epochs) where FedProx particularly outperforms FedAvg?
-

EXERCISE 5 - Client Sampling

Objective: Implement two client sampling strategies from the research paper ["On the Convergence of FedAvg on Non-IID Data" \(https://arxiv.org/abs/1907.02189\)](https://arxiv.org/abs/1907.02189).

EXERCISE 5.1 - Uniform Sampling Without Replacement

Background

Understand uniform sampling as described in Assumption 6. This involves selecting a subset of clients $|S_t| = K$ at each round without replacement. Understand the aggregation formula given by
$$\bar{w}_t \leftarrow \frac{1}{N} \sum_{k \in S_t} p_k w_k^t$$

Instructions

1. In `aggregator.py`, complete the `sample_clients()` method to uniformly sample `self.n_clients_per_round` clients from the total available clients.
2. Use `self.rng.sample` to sample `self.n_clients_per_round` unique ids from a population ranging from 0 to `self.n_clients - 1`.
3. Assign the list of sampled ids to `self.sampled_clients_ids`.
4. Modify the `mix()` method to:
 - Use only the sampled clients for training. For local training, loop over `self.sampled_clients_ids` instead of all clients.
 - Aggregate updates from the sampled clients. Adjust weights accordingly.

Run the code

Run the `train.py` script with `sampling_rate = 0.2`.

EXERCISE 5.2 - Sampling With Replacement

Background

Understand sampling with replacement according to sampling probabilities p_1, \dots, p_N . The aggregation formula adjusts to $w_t \leftarrow \frac{1}{K} \sum_{k \in S_t} w^k_t$.

Instructions

1. Extend the `sample_clients()` method to support sampling with replacement based on `self.sample_with_replacement` flag.
2. If `self.sample_with_replacement` is `True`, use `self.rng.choices` to sample clients based on their weights `self.clients_weights`.

Run the code

Run the `train.py` script with `sampling_rate = 0.2` and `sample_with_replacement = True`.

Good luck, and don't hesitate to ask questions and collaborate with your peers!

At the end of the lesson, you can send your document and code to: francesco.diana@inria.fr
(<mailto:francesco.diana@inria.fr>)