

Federated Learning & Data Privacy, 2024-2025

Nguyen Dinh Huy

Third Lab - 11 February 2025

Welcome to the third lab session of the Federated Learning & Data Privacy course! Today we will see how to implement Federated Learning in real network systems.

I. EXERCISE 5 - Federated Learning with Flower

Objective: Gain practical experience with the [Flower federated learning framework](#) by deploying a real, distributed federated learning experiment. Explore personalized federated learning algorithms and compare their performance against the standard FedAvg.

II. EXERCISE 5.1 - Federated Learning on Real Networks

Goal: Deploy a federated learning system using PyTorch and Flower to understand the setup and execution of federated learning in a networked environment.

Setup

1. **Clone the Flower repository and set up the PyTorch quickstart example:**

```
git clone --depth=1 https://github.com/adap/flower.git && mv
flower/examples/quickstart-pytorch . && rm -rf flower && cd quickstart-
pytorch
```

2. **Install the dependencies:**

```
pip install -e .
```

Run Federated Learning

- **Launch the simulation**

```
flwr run .
```

- Observe the federated training process initiated by PyTorch through Flower.

You can find complete instructions [here](#).

III. EXERCISE 5.2 - Tackling Data Heterogeneity with FedProx

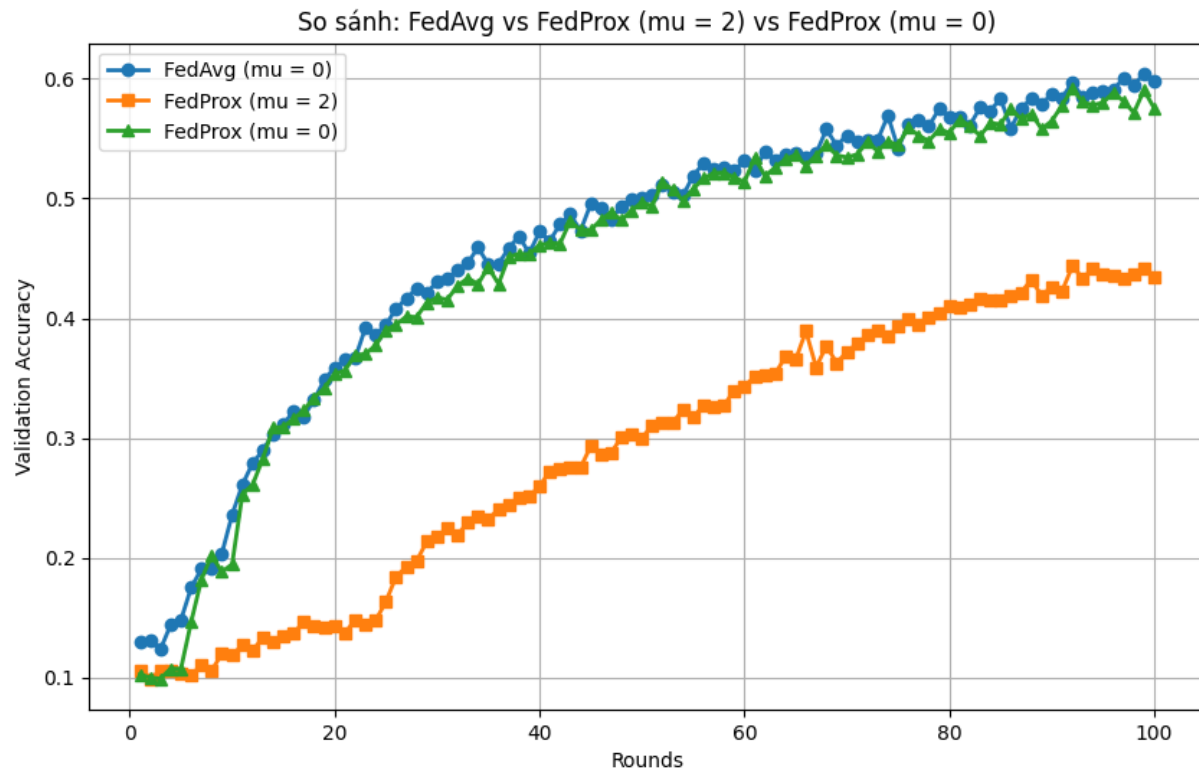
Objective: Understand how the FedProx algorithm addresses the challenges posed by data heterogeneity in federated learning and compare its performance with the FedAvg algorithm.

- **FedProx Overview:** FedProx modifies the local training objective by introducing a proximal term, which aims to reduce local model drift by penalizing significant deviations from the global model. Review the FedProx algorithm [Federated Optimization in Heterogeneous Networks \(Algorithm 2\)](#).
- **Instructions:** Follow the tutorial on FedProx available at [Flower's documentation](#). Run experiments trying different number of rounds and different μ values.

- **Analysis:** Discuss the observed differences in performance between FedAvg and FedProx. Are there specific configurations (e.g., number of local epochs) where FedProx particularly outperforms FedAvg?

RESULT :

Code Link : <https://github.com/DinhHuy1405/Federated-Learning-2024---2025/tree/main/TP3>



Observations from the Experiment

The graph shows validation accuracy over **100 rounds** for three federated learning approaches:

1. **FedAvg ($\mu = 0$)** → Traditional federated averaging.
2. **FedProx ($\mu = 2$)** → FedProx with a strong regularization term.
3. **FedProx ($\mu = 0$)** → Equivalent to FedAvg.

Key Insights:

1. FedAvg ($\mu = 0$) vs. FedProx ($\mu = 0$):

- As expected, **FedAvg ($\mu = 0$)** and **FedProx ($\mu = 0$)** behave similarly because FedProx with $\mu = 0$ is equivalent to FedAvg.
- There are slight performance variations, but both approaches show almost identical trends in accuracy improvement over time.
- Any small differences could be attributed to implementation details or stochastic effects.

2. FedAvg ($\mu = 0$) vs. FedProx ($\mu = 2$):

- **FedAvg ($\mu = 0$)** achieves significantly **higher accuracy** over time compared to **FedProx ($\mu = 2$)**.
- **FedProx ($\mu = 2$)** starts with a similar accuracy but **converges much more slowly**.
- This suggests that a strong proximal term ($\mu = 2$) effectively **reduces client drift** but also **restricts adaptation**, leading to **lower overall accuracy**.

3. FedProx ($\mu = 2$) vs. FedProx ($\mu = 0$):

- **FedProx ($\mu = 0$)** consistently **outperforms FedProx ($\mu = 2$)**.
- This indicates that for this dataset and training setup, a **strong proximal constraint ($\mu = 2$)** **overly restricts local model updates**, preventing optimal learning.

Analysis of Data Heterogeneity

- **FedProx** is designed to mitigate **data heterogeneity** by limiting client model divergence.
- However, selecting an appropriate μ value is critical:
 - **Small μ (close to 0)**: Allows more flexibility but may suffer from client drift.
 - **Large μ** : Reduces drift but slows down convergence, leading to suboptimal performance.

Impact of Local Epochs

- The impact of FedProx **may vary depending on the number of local epochs**:
 - **If local training lasts multiple epochs**, FedAvg may diverge more due to heterogeneous data distributions.
 - **FedProx can help by constraining local updates**, keeping them closer to the global model.
 - **In scenarios with higher local epochs**, FedProx might perform better than FedAvg.

IV. BONUS EXERCISE - Personalized Federated Learning

Goal: Evaluate a personalized federated learning algorithm using Flower, showing its possible advantages over the FedAvg algorithm.

Choose one of the two proposed personalization algorithms:

1. Federated Learning with Personalization Layers (FedPer)


- **Overview:** FedPer implements personalization by allowing some neural network layers to be client-specific, making the model closer to individual data distributions.
- **Instructions:** Follow the tutorial on FedPer available at [Flower's documentation](#).

Evaluation

- Reproduce the tutorial and compare the results with FedAvg to highlight the benefits of personalization.

RESULT:

Code Link: <https://github.com/DinhHuy1405/Federated-Learning-2024---2025/tree/main/TP3>



Observations from the Experiment

The graph displays the validation accuracy over **100 rounds** for three federated learning strategies:

1. **FedAvg ($\mu = 0$)** → Traditional federated averaging.
2. **FedProx ($\mu = 2$)** → FedProx with a strong regularization term.
3. **FedPer** → An approach focusing on personalized federated learning.

Key Insights:

1. FedAvg ($\mu = 0$):

- **FedAvg ($\mu = 0$)** shows the highest and most stable performance, achieving over **60% accuracy**.
- It is the most effective strategy among the three, displaying rapid and consistent improvements in accuracy.

2. FedProx ($\mu = 2$):

- **FedProx ($\mu = 2$)** starts with lower accuracy but gradually improves, stabilizing around **40% accuracy**.
- The strong regularization parameter ($\mu = 2$) seems to limit performance initially but ensures moderate improvements over time.

3. FedPer:

- **FedPer**, aiming at personalized models, shows much lower and fluctuating performance, with accuracy hovering around **10%**.
- This suggests potential issues in personalization effectiveness or instability in the learning process under this specific setup.

Comparison and Analysis:

- **FedAvg ($\mu = 0$) vs. FedProx ($\mu = 2$):**

 - As seen before, **FedAvg** consistently outperforms **FedProx with a strong μ** . The lack of a regularization constraint allows for more adaptable but possibly more varied client updates.
 - **FedProx** reduces client drift but at a significant cost to convergence speed and ultimate performance.

- **FedAvg ($\mu = 0$) vs. FedPer:**

 - **FedAvg** significantly outperforms **FedPer**, indicating that the non-personalized model is more effective in this scenario.
 - **FedPer's** performance suggests that its approach to personalization might not be effectively capturing useful patterns or it might be too sensitive to client-specific noise.

- **FedProx ($\mu = 2$) vs. FedPer:**

 - Even with its conservative updates, **FedProx** performs better than **FedPer**, further highlighting the challenges in the personalization strategy employed by **FedPer**.